

Introduction to Digital
Image Processing with
MATLAB® Asia Edition
McAndrew-Wang-Tseng

Chapter 1: Introduction

CENGAGE
Learning™

1.2 What Is Image Processing?

- **Image processing** involves changing the nature of an image in order to either
 1. improve its pictorial information **for human interpretation** (**enhancement**) → Category 1
 2. render it more suitable **for autonomous machine perception** (**recognition**) → Category 2

CENGAGE
Learning™

1.2 What Is Image Processing?

Category 1.

- Enhancing the edges of an image to make it appear sharper
- Note how the second image appears cleaner; it is a more pleasant image.
- Sharpening edges is a vital component of printing.



FIGURE 1.1 Image sharpening. (a) The original image. (b) Result after sharpening.

CENGAGE
Learning™

1.2 What Is Image Processing?

- Removing noise from an image
- Noise in the image comes from random errors.

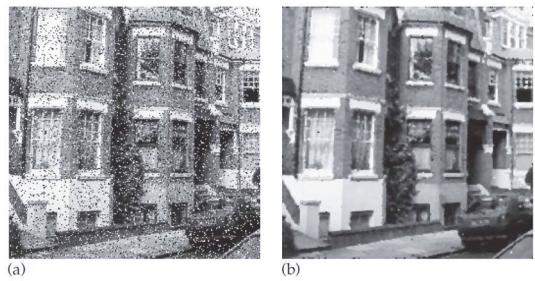


FIGURE 1.2 Removing noise from an image. (a) The original image. (b) After removing noise.

CENGAGE
Learning™

1.2 What Is Image Processing?

- Removing motion blur from an image. An example is given in Figure 1.3
- In Figure 1.3(b), it is easier to read the number plate and to see the spikes on the fence behind the car, as well as other details not at all clear in the original image Figure 1.3(a).



FIGURE 1.3 Image deblurring. (a) The original image. (b) After removing the blur.

CENGAGE
Learning™

1.2 What Is Image Processing? Category 2

- Obtaining the edges of an image
- Once we have the edges we can measure their spread and the area contained within them.
- We can also use edge-detection algorithms as a first step in edge enhancement

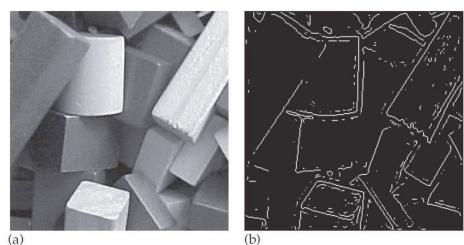


FIGURE 1.4 Finding edges in an image. (a) The original image. (b) Its edge image.

CENGAGE
Learning™

1.2 What Is Image Processing?

- For measurement or counting purposes, we may not be interested in all the detail in an image (Figure 1.5).
- We could, for example, measure the size and shape of the animal without being distracted by unnecessary detail.

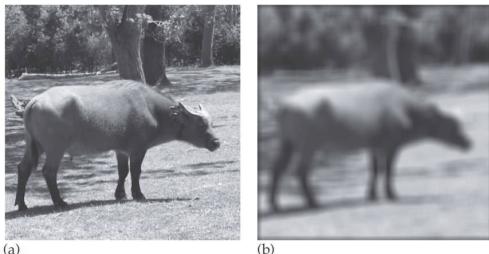


FIGURE 1.5 Blurring an image. (a) The original image. (b) Blurring to remove detail.

CENGAGE Learning

Example of Measuring Size and Shape by Noise Reduction

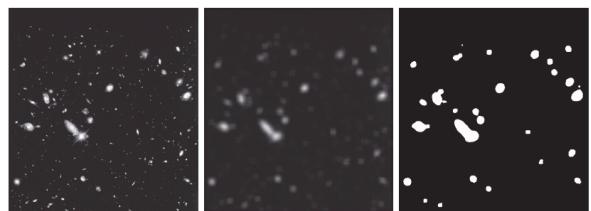


FIGURE 3.36 (a) Image from the Hubble Space Telescope. (b) Image processed by a 15×15 averaging mask. (c) Result of thresholding (b). (Original image courtesy of NASA.)

CENGAGE Learning

1.3 Image Sampling and Acquisition

- Sampling** refers to the process of digitizing a continuous function
 - e.g., $y = \sin(x) + \frac{1}{3} \sin(3x)$
 - ✓ sample it at **10** evenly spaced values of x only
 - ✓ sample it at **100** points



FIGURE 1.6 Sampling a function—undersampling.

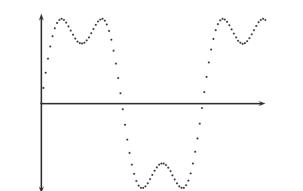


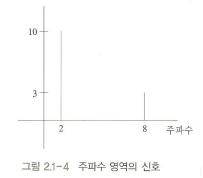
FIGURE 1.7 Sampling a function with more points.

CENGAGE Learning

How many samples ?

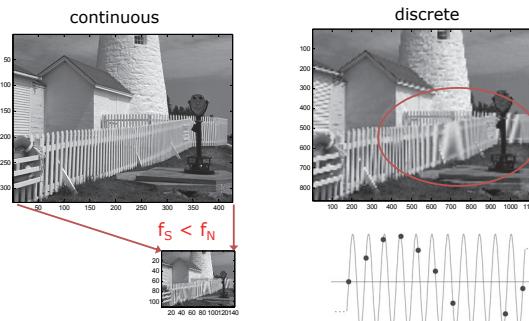
- Sampling rate :** samples/sec
- Nyquist sampling theorem (= Shannon's sampling theorem)**
 - gives us an **optimal number of samples** which also indicates optimal trade-offs between memory space and loss of information.
 - If a continuous signal is band-limited, i.e. if a signal has a maximum frequency f_{max} , we have to choose the sampling frequency $f_s > 2f_{max}$ \rightarrow Nyquist rate
- Examples**
 - Find the Nyquist rate of the signal ?
 $8 \times 2 = 16$
means what? We need at least 16 Hz of sampling frequency
 - Human ears can hear the signal only within the range of 20~20KHz \rightarrow sampling rate of compact disk : 44.1KHz $>$ 40KHz
 - Bandwidth of conventional telephone line: 0.3~3.4KHz
 \rightarrow sampling rate : 8KHz $>$ 6.8 KHz

Nyquist frequency



8

Aliasing Example



- Sampling below the Nyquist rate
 \Rightarrow causes distortion in restored signal.
(Especially high frequency component)

CENGAGE Learning

Aliasing Example



FIGURE 4.17 Illustration of aliasing on resampled images. (a) A digital image with negligible visual aliasing. (b) Result of resizing the image to 50% of its original size by pixel deletion. Aliasing is clearly visible. (c) Result of blurring the image in (a) with a 3×3 averaging filter prior to resizing. The image is slightly more blurred than (b), but aliasing is not longer objectionable. (Original image courtesy of the Signal Compression Laboratory, University of California, Santa Barbara.)

CENGAGE Learning

1.3 Image Sampling and Acquisition

- **Shannon's Sampling Theorem** which says, in effect, that a continuous function can be reconstructed from its samples provided that the sampling frequency is at least twice the maximum frequency(Nyquist frequency) in the function.

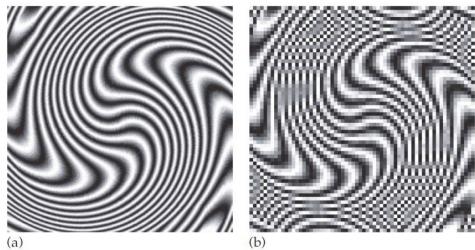


FIGURE 1.8 Effects of sampling. (a) Correct sampling; no aliasing. (b) An undersampled version with aliasing.

CENGAGE Learning

Image Acquisition

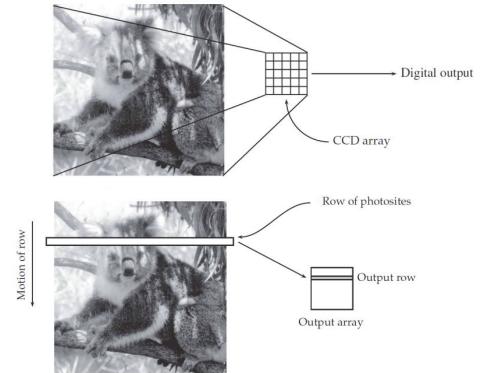


FIGURE 1.10 Capturing an image with a CCD scanner.

CENGAGE Learning

ENERGY SOURCES FOR MEDICAL IMAGING

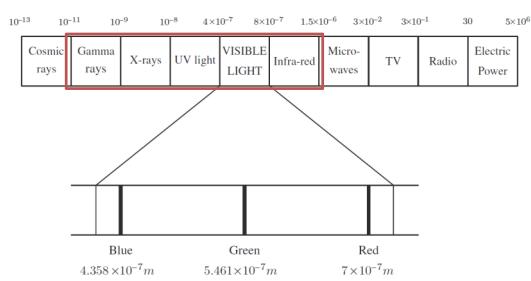


FIGURE 1.11 The electromagnetic spectrum.

CENGAGE Learning

Image Acquisition by X-ray Computed Tomography

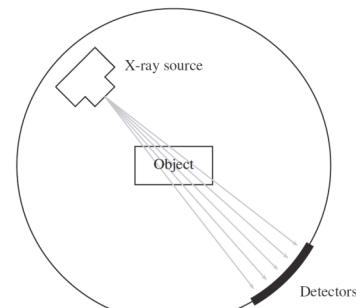


FIGURE 1.12 X-ray tomography.

CENGAGE Learning

1.4 Images and Digital Images

- We may consider this image as being a two-dimensional function $f(x, y)$
- We may assume that in such an image, brightness values can be any real numbers in the range **0.0 (black)** to **1.0 (white)**
- The $f(x, y)$ values in a digital image take only integer values ranging from **0 (black)** to **255 (white)**

CENGAGE Learning

Pixel Values in an Image

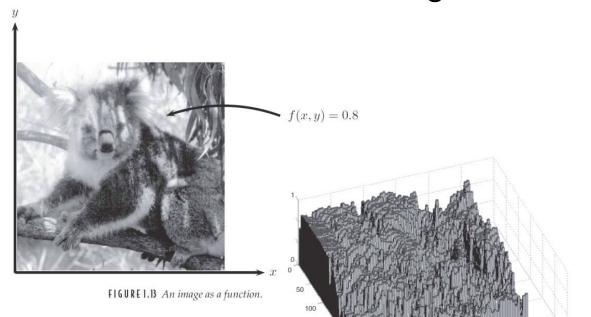


FIGURE 1.13 An image as a function.

FIGURE 1.14 The image of Figure 1.13 plotted as a function of two variables.

1.4 Images and Digital Images

- A **digital image** can be considered as a large array of sampled points from the continuous image.
- These points are the **pixels**, which constitute the digital image.
- The pixels surrounding a given pixel constitute its **neighborhood**.

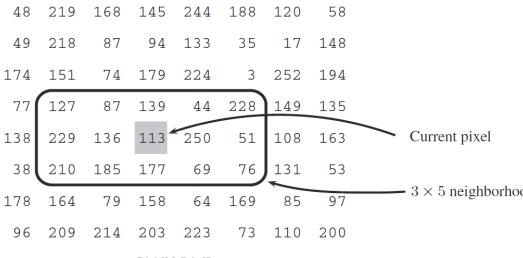


FIGURE 1.15 Pixels with a neighborhood.

CENGAGE Learning

Applications of Image Processing

- IMAGE ENHANCEMENT: Category 1**
 - ✓ sharpening or deblurring an out-of-focus image
 - ✓ highlighting edges,
 - ✓ improving image contrast or brightening an image, and
 - ✓ removing noise
- IMAGE RESTORATION: Category 1**
 - ✓ removing of blur caused by linear motion,
 - ✓ removal of optical distortions, and
 - ✓ removing periodic interference
- IMAGE SEGMENTATION: Category 2**
 - ✓ finding lines, circles, or particular shapes in an image, and
 - ✓ identifying cars, trees, buildings, or roads in an aerial photograph
- A given algorithm may be used for both image enhancement or for image restoration

CENGAGE Learning

Type of Digital Image

• Binary

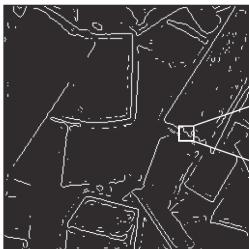


FIGURE 1.16 A binary image.

CENGAGE Learning

Type of Digital Image

• True color or red-green-blue(RGB)



Red Green Blue

FIGURE 1.18 A true color image.

CENGAGE Learning

Type of Digital Image

• Grayscale

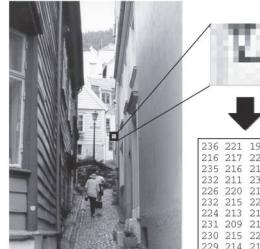


FIGURE 1.17 A grayscale image.

CENGAGE Learning

Type of Digital Image

• Indexed Color

- Most color images contain only a small subset of 16 million possible colors in RGB.
- For convenience of storage and file handling, we might need a list of all colors used in the image. => **colormap**



Indices	Color map
4	0.1211 0.1211 0.1416
5	0.1807 0.2549 0.1729
4	0.2197 0.3447 0.1807
5	0.1611 0.1768 0.1924
5	0.2432 0.2471 0.1924
5	0.2119 0.1963 0.2002
5	0.2627 0.2588 0.2549
5	0.2197 0.2432 0.2588
6	⋮ ⋮ ⋮
6	11 20 33 33 58 37

FIGURE 1.19 An indexed color image.

CENGAGE Learning

1.9 Image File Sizes

- A **512 × 512 binary image**

$$\begin{aligned} 512 \times 512 \times 1 &= 262,144 \text{ bytes} \\ &= 32768 \text{ bytes} \\ &= 32.768 \text{ Kb} \\ &\approx 0.033 \text{ Mb} \end{aligned}$$

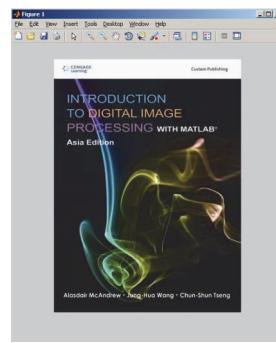
- A **grayscale image**

$$\begin{aligned} 512 \times 512 \times 1 &= 262,144 \text{ bytes} \\ &= 262.14 \text{ Kb} \\ &\approx 0.262 \text{ Mb.} \end{aligned}$$

- A **RGB color image**

$$\begin{aligned} 512 \times 512 \times 3 &= 786,432 \text{ bytes} \\ &= 786.43 \text{ Kb} \\ &\approx 0.786 \text{ Mb.} \end{aligned}$$

 CENGAGE
Learning



Chapter 2: Images and MATLAB

 CENGAGE
Learning

Matlab Image Processing Toolbox

- MATLAB is a data analysis software package with powerful support for matrices and matrix operations

✓ Command window

>>

✓ Reads pixel values from an image

>> w=imread('wombats.tif');

 CENGAGE
Learning

How to Display Grayscale Images

>> figure, imshow(w)

figure

- ✓ Creates a figure on the screen
- ✓ A figure is a window in which a graphics object can be placed

imshow (w)

- ✓ displays the matrix w as an image

 CENGAGE
Learning

Pixval Command

pixval on

- ✓ turns on the pixel values in our figure.
- ✓ They appear at the bottom of the figure in the form.

c, r = p



FIGURE 2.1 The wombats image with pixval on.

Note that the command `pixval on` may be removed in a later MATLAB version

 CENGAGE
Learning

2.2 RGB Images

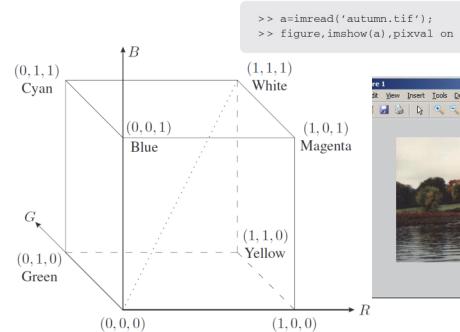


FIGURE 2.2 The color cube for the RGB color model.



 CENGAGE
Learning

2.2 RGB Images

- Multidimensional array

```
>> size(a) % 206 (rows) 345 (columns) 3 (pages)
>> a(100,200,2) % 25
>> a(100,200,1:3) % 75 25 30
>> a(100,200,:) % 75 25 30
>> impixel(a,200,100) % 75 25 30
```



2.3 Indexed Color Images

```
>> figure, imshow('trees.tif')
>> em = imread('trees.tif');
>> figure, imshow(em)
```



```
>> [em, emap] = imread('trees.tif');
>> figure, imshow(em, emap)
```



How to Know Information on Your Image

- A great deal of information can be obtained with the `imfinfo` function (pp. 26)
- Useful header information
 - Filename, FileSize, Format: 'tif'
 - Width, Height, BitDepth
 - ColorType: 'indexed', 'truecolor', 'grayscale'
 - ByteOrder: 'little-endian', 'big-endian'
 - Compression: compression algorithm used for the image
 - XResolution, YResolution, ResolutionUnit: 'Inch'
 - Colormap: [256x3 double] ← 'indexed'
- `imread` automatically read the header of the specified image file.



2.4 Data Types and Conversions

```
>> a=23;
>> b=uint8(a);
>> b
b =
23
>> whos a b
  Name      Size            Bytes  Class
  a            1x1             8  double array
  b            1x1             1  uint8 array
```

- You can use MATLAB for image processing without ever really knowing the difference among GIF, TIFF, PNG, etc.
- However, some knowledge of the different **graphics formats** (**data type**) can be useful in order to make a reasoned decision.



2.4 Data Types and Conversions

TABLE 2.1 Data types in MATLAB

Data Type	Description	Range
int8	8-bit integer	-128-127
uint8	8-bit unsigned integer	0-255
int16	16-bit integer	-32768-32767
uint16	16-bit unsigned integer	0-65535
double	Double precision real number	Machine specific

TABLE 2.2 Converting images in MATLAB

Function	Use	Format
ind2gray	Indexed to grayscale	y=ind2gray(x,map);
gray2ind	Grayscale to indexed	[y,map]=gray2ind(x);
rgb2gray	RGB to grayscale	y=rgb2gray(x);
gray2rgb	Grayscale to RGB	y=gray2rgb(x);
rgb2ind	RGB to indexed	[y,map]=rgb2ind(x,n);
ind2rgb	Indexed to RGB	y=ind2rgb(x,map);



2.5 Images Files and Formats

- The `imread` and `imwrite` functions of MATLAB currently support the following formats.
 - ✓ JPEG: These images are created using the Joint Photographic Experts Group compression method (Ch14).
 - ✓ TIFF: A very general format that supports different compression methods, multiple images per file, and binary, grayscale, truecolor, and indexed images. Tagged Image File Format.
 - ✓ GIF: A format designed for data transfer. It is still popular and well supported, but is somewhat restricted in the image types it can handle. Graphics Interchange Format.
 - ✓ BMP: Microsoft Bitmap format has become very popular and is used by Microsoft operating systems.
 - ✓ PNG, HDF, PCX, XWD, ICO, CUR



HEXADECIMAL DUMP

- We can see inside of image file as it is in binary.

```
function dumphex(filename, n)
%
% DUMPHEX(FILENAME,N) prints the first 16*n bytes of the file FILENAME
% in hex and ASCII. For example:
%
%   % dumphex('picture.bmp',4)
%
fid = fopen(filename, 'r');
if fid==1
    error('File does not exist or is not in your Matlab path');
end;
a=fread(fid,16*n,'uchar');
idx=find(a>=32 & a<=126); % check if it is an ascii code
ah=dec2hex(a);
b=repmat([' '],16*n,3);
b2=repmat([' '],16,n);
b2(idx)=char(a(idx));
b(:,1:2)=ah;
[reshape(b',48,n)' repmat(' ',n,2) reshape(b2,16,n)']
```

FIGURE 2.3 A function for producing hexadecimal dumps.

CENGAGE
Learning

Vector versus Raster Images

• Vector images

- A collection of **lines or vectors**
- Can be magnified to any desired size without losing sharpness
- Good for** images consisting mostly of **lines** and mathematically described **curves (characters, graphs)**
- PostScript by Adobe Acrobat

• Raster images

- A collection of **dots**
- A list of gray or color intensities of each pixel
- Images** captured by digital camera, scanner, etc.
- The great bulk of image file formats store images as raster information.

CENGAGE
Learning

2.5.2 A Simple Raster Format

- As well as containing all pixel information, an image file must contain some **header information**
 - This must include the size of the image, but may also include some documentation, a color map, and the compression method.
- Example: PGM format
 - designed to be a generic format used for conversion between other formats

```
P2 means ASCII PGM file
# CREATOR: The GIMP's PNM Filter Version 1.0 comments
256 256 # of cols and rows
255 # of grayscales
41 53 53 53 53 49 49 53 53 56 56 49 41 46 53 53 53
53 41 46 56 56 56 53 53 46 53 41 41 53 56 49 39 46
```

pixel values delimited by space and CR

CENGAGE
Learning

Header Format of Microsoft BMP

Bytes	Information	Description
0~1	Signature	BM in ASCII = [42 4D] in hexadecimal.
2~5	FileSize	The size of the file in bytes.
6~9	Reserved	All zeros.
10~13	DataOffset	File offset to the raster data.
14~17	Size	Size of the information header = 40 bytes.
18~21	Width	Width of the image in pixels.
22~25	Height	Height of the image in pixels.
26~27	Planes	Number of image planes (= 1).
28~29	BitCount	Number of bits per pixel: <ul style="list-style-type: none"> [1: Binary images; two colors, 4: $2^4 = 16$ colors (indexed), 8: $2^8 = 256$ colors (indexed), 16: 16-bit RGB; $2^{16} = 65,536$ colors, 24: 24-bit RGB; $2^{24} = 17,222,216$ colors.

CENGAGE
Learning

Header Format of Microsoft BMP

Bytes	Information	Description
30~33	Compression	Type of compression used: <ul style="list-style-type: none"> 0: No compression (most common), 1: 8-bit RLE encoding (rarely used), 2: 4-bit RLE encoding (rarely used).
34~37	ImageSize	Size of the image. If compression is 0, then this value may be 0.
38~41	HorizontalRes	The horizontal resolution in pixels per meter.
42~45	VerticalRes	The vertical resolution in pixels per meter.
46~49	ColorsUsed	The number of colors used in the image. If this value is zero, then the number of colors is the maximum obtainable with the bits per pixel, that is, 2^{BitCount} .
50~53	ImportantColors	The number of important colors in the image. If all the colors are important, then this value is set to zero.

CENGAGE
Learning

Reading Header of Microsoft BMP

```
>> dumphex('blocksets.bmp', 4)
ans =
42 4D 6E 18 00 00 00 00 00 36 00 00 00 28 00 BM.....6...
00 00 [42 00 00 00] 1F 00 00 01 00 18 00 00 ..B.....
00 00 38 18 00 00 C4 02 00 00 C4 0E 00 00 00 00 ..8....
00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 .....
```

- The image width is given by bytes 18~21; they are in the second row
42 00 00 00
- Since BMP uses '**little-endian**' **for byte ordering**, to find the actual width, we reorder these bytes back-to-front:
00 00 00 42
- Now we can convert to decimal $(4 \times 16^1) + (2 \times 16^0) = 66$ which is the image width in pixels
- The image height
1F 00 00 00
 $(1 \times 16^1) + (F \times 16^0) = 31$

CENGAGE
Learning

GIF [jif]

- Proposed by Compuserve in the late 1980's as a mean for distributing images over networks.
- Colors are stored using a **color map**. The GIF specification allows a **maximum of 256 colors** per image
- GIF doesn't allow binary or grayscale images, except as can be produced with RGB values
- The pixel data is compressed using **LZW** (Lempel-Ziv-Welch) compression
- The GIF format allows **multiple images per file**. This aspect can be used to create **animated GIFs**



Portable Network Graphics(PNG[ping])

- The PNG format has been more recently designed to replace GIF and to overcome some of GIF's disadvantages.
- Does not rely on any patented algorithms, and it supports more image types than GIF.
- Supports grayscale, **true color**, and indexed images.
- Moreover, its compression utility, **zlib**, always results in genuine compression



JPEG

- The JPEG algorithm uses **lossy compression**, in which not all the original data can be recovered

```
>> dumphex('ngc6543a.jpg',4)
ans =
FF D8 FF E0 00 10 4A 46 49 46 00 01 01 00 00 01 .....JFIF.....
00 01 00 00 FF FE 00 47 43 52 45 41 54 4F 52 3A .....GCREATOR:
20 58 56 20 56 65 72 73 69 6F 6B 20 33 2E 30 30 XV Version 3.00
62 20 20 52 65 76 3A 20 36 2F 31 35 2F 39 34 20 b Rev: 6/15/94
```



JPEG Header

Bytes	Information	Description
0-1	Start of image marker	Always FF D8.
2-3	Application marker	Always FF E0.
4-5	Length of segment	
6-10	JFIF\ 0	ASCII JFIF.
11-12	JFIF version	In our example above 01 01 or version 1.1. Values are: 0 arbitrary units; 1 pixel/inch; 2 pixels/centimeter.
13	Units	
14-15	Horizontal pixel density	
16-17	Vertical pixel density	
18	Thumbnail width	If this is 0, there is no thumbnail.
19	Thumbnail height	If this is 0, there is no thumbnail.



TIFF

- One of **the most comprehensive(complete)** image formats
- Can store **multiple images per file**
- Allows **different compression routines** and **different byte orderings** (none, LZW, JPEG, Huffman, run-length, etc)
- Allows binary, grayscale, truecolor or indexed images, and opacity or transparency
- An excellent format for data exchange
- Almost all journal papers accept tiff files** as their standard figure format.



TIFF

Bytes	Information	Description
0-1	Byte order	Either 4D 4D: ASCII MM for big endian, or 49 49: ASCII II for little endian.
2-3	TIFF version	Always 00 2A or 2A 00 (depending on the byte order) = 42.
4-8	Image offset	Pointer to the position in the file of the data for the first image.

```
>> dumphex('newborn.tif',4)
ans =
49 49 2A 00 E0 01 01 00 32 7C 5B 2D 23 19 0B 15 II*....2|[-#...
10 0E 0D 0F 10 0F 0E 10 11 11 0E 12 13 12 10 17 .....p.....
10 1D 70 8E 99 A0 AE B5 BB BA C2 C6 C6 CB D3 D0 D2 D1 CA DB DE E1 E5 E6 DF E4 E9 FE EB 0B ED .....
```



TIFF

- This particular image uses the little-endian byte ordering.
- The first image in this file (which is in fact the only image), begins at byte

```
E0 01 01 00
```
- Because this is a **little-endian** file, we **reverse the order** of the bytes: **00 01 01 E0**. This works out to **66016**.

big endian ordering: E0 at the lowest address(say 1000), 00 at the highest address (then 1003)



DICOM

- DICOM (Digital Imaging and Communications in Medicine)
- Like GIF, may hold multiple image files.
- May be considered as slices or frames of a three dimensional object.
- The DICOM specification is huge and complex. Drafts have been published on the World Wide Web.
- Header contains **protocol parameters of image modalities**(CT, MR, PET, SPECT, X-ray etc) as well as image information itself.



Storing an Image into a File in MATLAB

```
imwrite(X,map,'filename','fmt')
```

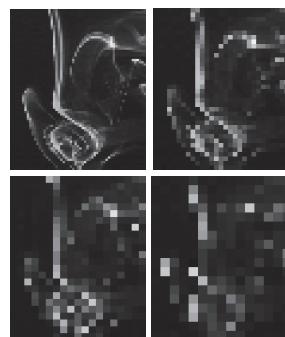
- Which writes the image stored in matrix **X** with color map **map** (if appropriate) to file **filename** with format **fmt**.
- **Without the map argument**, the image data is supposed to be **grayscale or RGB**.

e.g.

```
a=imread('autumn.tif');
```



```
imwrite(a,'autumn.png','png');
```



Chapter 3: Image Display



Introduction

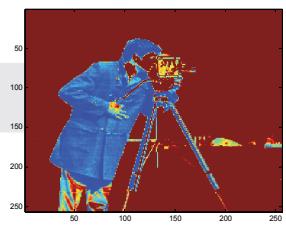
- We take a **closer look at** the use of the **imshow** function.
- We look at **image quality** and how that may be affected by various image attributes.
- For human vision in general, the preference is for images to be sharp and detailed.
- There are many factors that will affect the display
 - ✓ ambient lighting,
 - ✓ the monitor type and settings,
 - ✓ the graphics card, and
 - ✓ monitor resolution



3.2 Basics of Image Display

- This function **image** simply displays a matrix as an indexed color image with default color map named **jet** that consists of 64 bright colors.

```
>> c=imread('cameraman.tif');  
>> image(c)
```



3.2 Basics of Image Display

- To display the image properly, we need to add several **extra commands** to the **image** line

```
>> size(unique(c))
ans =
247    1

>> image(c), truesize, axis off, colormap(gray(247))

By default, ✓ turns off
display axis      adjust color map to
256x256 image labeling use shades of gray
only
```



GRAY(M) returns an M-by-3 matrix containing a **gray-scale colormap**. GRAY, by itself, is the same length as the current figure's colormap.

CENGAGE Learning

3.2 Basics of Image Display

- We may to adjust the color map to use fewer or more colors; however, this can have a dramatic effect on the result.

colormap(gray(512)) colormap(gray(128))



CENGAGE Learning

3.2 Basics of Image Display

- use **imread** to pick up the color map

```
>> [x, map]=imread('trees.tif');
>> image(x), truesize, axis off, colormap(map)
```

- ✓ **map** is <256x3 double> in the **workspace**



CENGAGE Learning

3.2 Basics of Image Display

- True color image will be read (by **imread**) as a **three-dimensional array**

- In such a case, **image** will ignore the current color map and assign colors to the display based on the values in the array

```
>> t=imread('twins.tif');
>> image(t), truesize, axis off
```



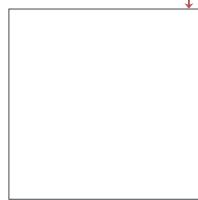
CENGAGE Learning

3.3 The imshow Function

- We have two choices with a matrix of type **double**:
 - ✓ Convert to type **uint8** and then display
 - ✓ Display the matrix directly
- imshow** will display a matrix of type **double** as a grayscale image (matrix elements are between 0 and 1)

```
>> c=imread('caribou.tif');
>> cd=double(c);
>> imshow(c), figure, imshow(cd)
```

imshow() assumes that min =0, max=1, then display values between 0 and 1 with the range of 0 to 255 intensity levels.



CENGAGE Learning

Possible Solutions ?

```
>> imshow(cd/512)
>> imshow(cd/128)
```



(a)



(b)

FIGURE 3.2 Scaling by dividing an image matrix by a scalar. (a) The matrix *cd* divided by 512. (b) The matrix *cd* divided by 128.

CENGAGE Learning

3.3 The imshow Function

- We can convert the original image to **double** more properly using the function **im2double**.

```
>> cd=im2double(c);
im2double() adjust all the pixel values so that min pixval = 0, max pixval = 1.
double() only changes data type, not pixel values.
```

If we take **cd** of type **double**, properly scaled so that all elements are between 0 and 1, we can convert it back to an image of type **uint8** in two ways:

```
>> c2=uint8(255*cd);
>> c3=im2uint8(cd);

im2uint8() adjust all the pixel values so that min pixval = 0, max pixval = 255. GAGE
Learning
```

More about imshow()

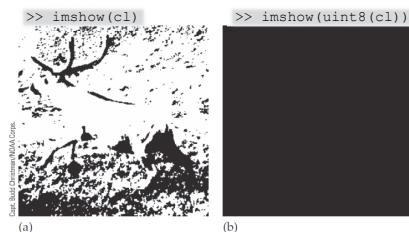


FIGURE 3.3 Making the image binary. (a) The caribou image turned binary. (b) After conversion to type **uint8**.
uint8() only changes the type, not a pixel value.
minus value → 0, bigger than 255 → 255.

A possible panacea : **figure, imshow(c, []);**

CENGAGE Learning

3.3 The imshow Function

- BINARY IMAGES:** MATLAB have a logical flag, where **uint8** values 0 and 1 can be interpreted as logical data

```
>> cl=c>120;
```

- Check **c1** with **whos**

Name	Size	Bytes	Class	Attributes
c1	256x256	65536	logical	

CENGAGE Learning

3.4 Bit-Planes

- It highlights the contributions made to the total image appearance by specific bits.

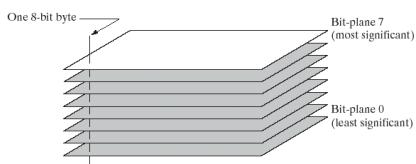


FIGURE 3.12
Bit-plane representation of an 8-bit image.

CENGAGE Learning

3.4 Bit Planes

- Grayscale images can be transformed into a sequence of binary images by breaking them up into their **bitplanes**
- The first bit plane(plane 0): LSB plane**
 - the least significant bit plane
- The eighth bit plane(plane 7): MSB plane**
 - the most significant bit plane

CENGAGE Learning

3.4 Bit Planes

- We start by making it a matrix of type **double**; this means we can perform arithmetic on the values

```
>> c=imread('cameraman.tif');
>> cd=double(c);
>> c0=mod(cd,2);
>> c1=mod(floor(cd/2),2);
>> c2=mod(floor(cd/4),2);
>> c3=mod(floor(cd/8),2);
>> c4=mod(floor(cd/16),2);
>> c5=mod(floor(cd/32),2);
>> c6=mod(floor(cd/64),2);
>> c7=mod(floor(cd/128),2);
```

CENGAGE Learning

FIGURE 3.4

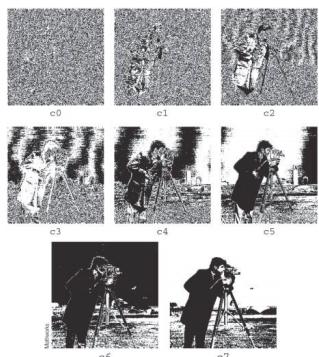


FIGURE 3.4 The bit planes of an 8-bit grayscale image.

CENGAGE Learning

Point Processing: Bit-Plane Slicing

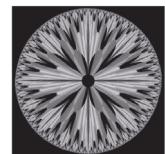


FIGURE 3.13 An 8-bit fractal image. (A fractal is an image generated from mathematical expressions.) (Courtesy of Ms. Melissa D. Blinde, Swarthmore College, Swarthmore, PA.)

- MSB 4bits에 시각적으로 중요한 데이터 집중, LSB 4bits는 세부묘사
- 영상 압축시 유용한 정보

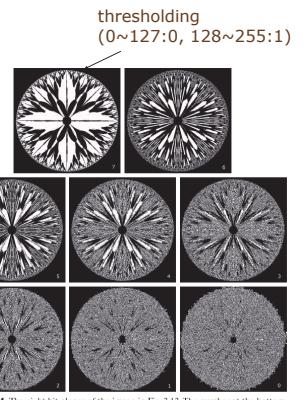


FIGURE 3.14 The eight bit planes of the image in Fig. 3.13. The number at the bottom right of each image identifies the bit plane.

CENGAGE Learning

Digital Image Representation & Image Resolution

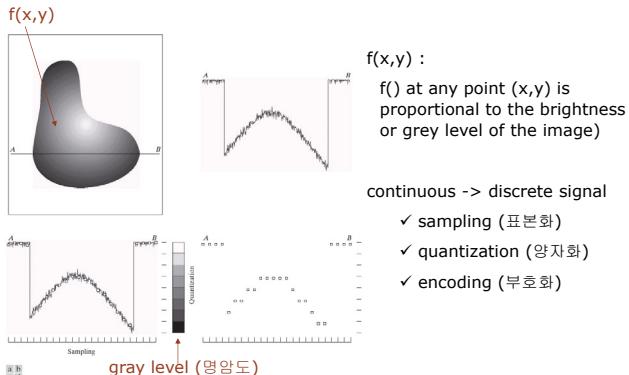
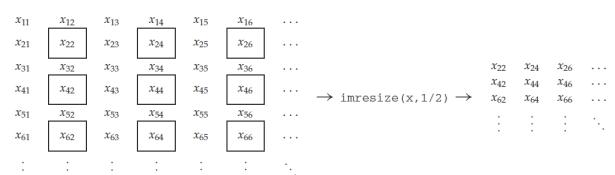


FIGURE 2.16 Generating a digital image. (a) Continuous image. (b) A scan line from A to B in the continuous image, used to illustrate the concepts of sampling and quantization. (c) Sampling and quantization. (d) Digital scan line.

CENGAGE Learning

3.5 Spatial Resolution

- The greater the spatial resolution, the more pixels are used to display the image.
- We can experiment with spatial resolution with MATLAB's **imresize** function



CENGAGE Learning

3.5 Spatial Resolution

Command	Effective resolution
<code>imresize(imresize(x, 1/4), 4);</code>	64×64
<code>imresize(imresize(x, 1/8), 8);</code>	32×32
<code>imresize(imresize(x, 1/16), 16);</code>	16×16
<code>imresize(imresize(x, 1/32), 32);</code>	8×8
<code>imresize(imresize(x, factor, 'nearest'), factor, 'nearest');</code>	

`imresize(imresize(x, factor, 'nearest'), factor, 'nearest');`

CENGAGE Learning

FIGURE 3.5



(a)



(b)

FIGURE 3.5 Reducing resolution of an image. (a) The original image. (b) Image at 128×128 resolution.

CENGAGE Learning

FIGURE 3.6

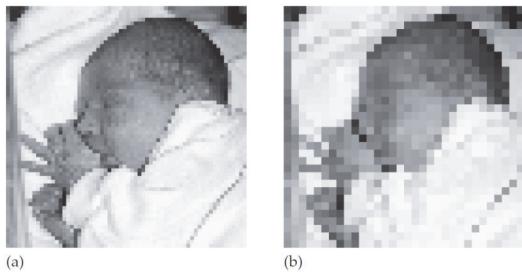


FIGURE 3.6 Further reducing the resolution of an image. (a) Image at 64×64 resolution. (b) Image at 32×32 resolution.

CENGAGE
Learning™

FIGURE 3.7

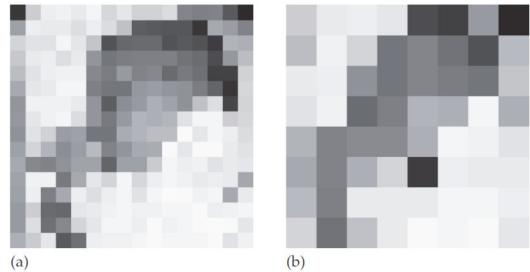


FIGURE 3.7 Even more reducing the resolution of an image. (a) Image at 16×16 resolution. (b) Image at 8×8 resolution.

CENGAGE
Learning™

3.6 Quantization

- Uniform quantization

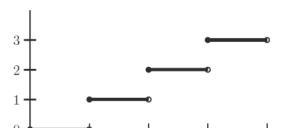


FIGURE 3.8 A mapping for uniform quantization.

Original values	Output value
0-63	0
64-127	1
128-191	2
192-255	3

CENGAGE
Learning™

3.6 Quantization

- To perform such a mapping in MATLAB, we can perform the following operations, supposing x to be a matrix of type `uint8`

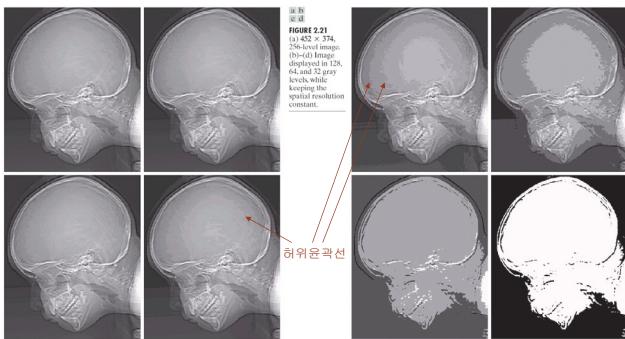
```
f=floor(double(x)/64);
q=uint8(f*64);
```

- There is a more elegant method of reducing the grayscales in an image, and it involves using the `grayslice` function

Command	Number of grayscales
<code>imshow(grayslice(x,128),gray(128))</code>	128
<code>imshow(grayslice(x,64),gray(64))</code>	64
<code>imshow(grayslice(x,32),gray(32))</code>	32
<code>imshow(grayslice(x,16),gray(16))</code>	16
<code>imshow(grayslice(x,8),gray(8))</code>	8
<code>imshow(grayslice(x,4),gray(4))</code>	4
<code>imshow(grayslice(x,2),gray(2))</code>	2

GRAYSICE creates indexed image from thresholding. CENGAGE Learning™

Quantization & Gray-level Resolution



- reduced gray level from 256 to 128,64,32,16,8,4,2
- 452x374 fixed spatial resolution
- best gray-level resolution in this example can be 64 => case by case

CENGAGE
Learning™

3.6 Dithering

- In general terms, it refers to the process of reducing the number of colors in an image.
- Representing an image with only two tones is also known as **halftoning**.

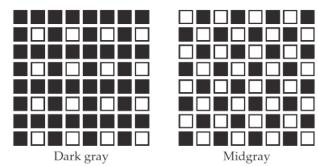


FIGURE 3.8 Patterns for dithering output.

CENGAGE
Learning™

Dithering

- A Standard dithering matrix

$$D = \begin{bmatrix} 0 & 128 \\ 192 & 64 \end{bmatrix} \quad D_2 = \begin{bmatrix} 0 & 128 & 32 & 160 \\ 192 & 64 & 224 & 96 \\ 48 & 176 & 16 & 144 \\ 240 & 112 & 208 & 80 \end{bmatrix}$$

- Suppose $d(i, j)$ is the matrix obtained by replicating the dithering matrix, then an output pixel $p(i, j)$ is defined by

$$p(i, j) = \begin{cases} 1 & \text{if } x(i, j) > d(i, j) \\ 0 & \text{if } x(i, j) \leq d(i, j) \end{cases}$$

 CENGAGE
Learning™

Dithering

- D or D_2 is repeated until it is as big as the image matrix, when the two are compared.

$$D = \begin{bmatrix} 0 & 128 \\ 192 & 64 \end{bmatrix}$$

$$D_2 = \begin{bmatrix} 0 & 128 & 32 & 160 \\ 192 & 64 & 224 & 96 \\ 48 & 176 & 16 & 144 \\ 240 & 112 & 208 & 80 \end{bmatrix}$$

```
>> D=[0 128;192 64];
>> r=repmat(D,128,128);
>> x2=x>r;imshow(x2)
>> D2=[0 128 32 160;192 64 224 96;48 176 16 144;240 112 208 80];
>> r2=repmat(D2,64,64);
>> x4=x>r2;imshow(x4)
```



 CENGAGE
Learning™

Dithering: Four Output Levels

- Dithering can be extended easily to more than two output gray values.
- For example, we wish to quantize to four output levels 0, 1, 2, and 3.

$$q(i, j) = [x(i, j)/85] \quad (\text{Since } 255/3 = 85)$$

$$p(i, j) = q(i, j) + \begin{cases} 1 & \text{if } x(i, j) - 85q(i, j) > d(i, j) \\ 0 & \text{if } x(i, j) - 85q(i, j) \leq d(i, j) \end{cases}$$

 CENGAGE
Learning™

Dithering : 4- and 8-Levels Grayscale

```
>> D = [0 56; 84 28];
>> r = repmat(D, 128, 128);
>> x = double(x);
>> q = floor(x/85);
>> x4 = q+(x-85*q>x);
>> imshow(uint8(85*x4))
```



(a)

```
>> D = [0 24; 36 12];
>> r = repmat(D, 128, 128);
>> x = double(x);
>> q = floor(x/37);
>> x8 = q+(x-37*q>x);
>> imshow(uint8(37*x8))
```



(b)

FIGURE 3.15 Dithering to more than two grayscales. (a) Dithering to four output grayscales. (b) Dithering to eight output grayscales.

 CENGAGE
Learning™

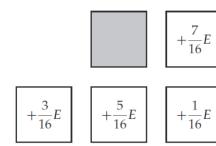
Error Diffusion

- The image is quantized at two levels.
- For each pixel we take into account the error between its gray value and its quantized value.
- The idea is to spread this error over neighboring pixels.

 CENGAGE
Learning™

Dithering: Floyd and Steinberg method

- Pixel values around 128 has big quantization error !
- For each pixel $p(i, j)$ in the image we perform the following sequence of steps:
 - Perform the quantization
 - Calculate the quantization error
 - Spread this error E over pixels to the right and below according to this table.



New $p(i,j) = E_{\text{old}}(p(i,j))$
: changes input pixels $p(i,j)$ before they are used.

 CENGAGE
Learning™

FIGURE 3.16

```

function y = floyd_Steinberg(x)
% FLOYD_STEINBERG applies Floyd-Steinberg error diffusion to an image x, which is
% assumed to be of type uint8.
%
height=size(x,1);
width=size(x,2);
y=int8(zeros(height,width));
z=zeros(height+2,width+2);
x(2:height+1,1:width+1)=x;
for i=2:height+1,
    for j=2:width+1,
        if x(i,j) < 128
            y(i-1,j-1) = 0;
            e = x(i,j);
        else
            y(i-1,j-1) = 255;
            e = x(i,j)-255;
        end
        z(i-1,j-1)=x(i-1,j-1)*7*e/16;
        z(i-1,j)=x(i-1,j)+1*e/16;
        z(i-1,j+1)=x(i-1,j+1)+5*e/16;
        z(i+1,j-1)=x(i+1,j-1)+9*e/16;
        z(i+1,j)=x(i+1,j)+4*e/16;
        z(i+1,j+1)=x(i+1,j+1)+e/16;
    end
end

```

FIGURE 3.16 A MATLAB function for applying Floyd-Steinberg error diffusion to a grayscale image.



FIGURE 3.17



FIGURE 3.17 The newborn baby image after Floyd-Steinberg error diffusion.



FIGURE 3.18

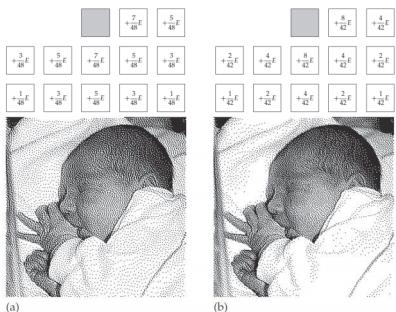


FIGURE 3.18 Using other error-diffusion schemes. (a) Result of Jarvis-Judice-Ninke error diffusion. (b) Result of Stucki error diffusion.



Summary

- **Chap 1. Introduction**

- ✓ Sampling, image pixels, image file size

- **Chap 2. Image and Matlab**

- ✓ Gray scale image, color images(RGB, indexed)
- ✓ Image file formats(GIF, TIFF, BMP, JPEG)
- ✓ Image header

- **Chap 3. Image Display**

- ✓ Bit planes
- ✓ Quantization
- ✓ Dithering
- ✓ Error diffusion

